Report: COVID-19 Data Analysis with R

Sandro Mikautadze

Introduction

This is exactly the TLDR you were looking for. **Enjoy!**

What is the project about?

The following project is an individual assignment taken from a course on Dataquest. The aim is to analyse a set of data collected from January 20th to June 1st 2020 taken from Kaggle. The research questions are the following:

- **Q1**: which countries have had the highest number of positive cases against the number of tests?

- **Q2**: which countries have made the best effort in terms of the number of COVID-19 tests conducted related to their population?

- **Q3**: which countries were ultimately the most and least affected related to their population?

The type of analysis is purely **descriptive**. All the snippets of code in this report use variables defined in the scripts.

Data

Raw data overview

We first analyse the dimensions, the column names and the information provided by each column in the data set:

```
glimpse(covid19_raw)
```

```
## Rows: 10,903
## Columns: 14
## $ Date                   <chr> "2020-01-20", "2020-01-22", "2020-01-22", "202~
## $ Continent_Name         <chr> "Asia", "North America", "North America", "Nor~
## $ Two_Letter_Country_Code <chr> "KR", "US", "US", "US", "US", "KR", "US", "US"~
## $ Country_Region         <chr> "South Korea", "United States", "United States~
## $ Province_State         <chr> "All States", "All States", "Washington", "All~
## $ positive               <int> 1, 1, 1, 1, 1, 2, 1, 1, 4, 0, 3, 0, 0, 0, 0, 1~
## $ hospitalized           <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
## $ recovered              <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
## $ death                  <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
```

```
## $ total_tested        <int> 4, 1, 1, 1, 1, 27, 1, 1, 0, 0, 0, 0, 0, 0, 0, ~
## $ active              <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
## $ hospitalizedCurr     <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
## $ daily_tested        <int> 0, 0, 0, 0, 0, 5, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
## $ daily_positive       <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
```

Looking at the data set repository from Kaggle we have further information on the meaning of each column:

- `Date`: date of the data collection

- `Country_Region`: country names

- `Province_State`: state/province names; value "All States" is put when state/provincial level data is NA

- `positive`: cumulative number of positive cases reported

- `active`: number of actively cases on that day

- `hospitalized`: cumulative number of hospitalized cases reported

- `hospitalizedCurr`: number of actively hospitalized cases on that day

- `recovered`: cumulative number of recovered cases reported

- `death`: cumulative number of deaths reported

- `total tested`: cumulative number of tests conducted

Note: Not all columns in our version of the data frame were present in the data description on Kaggle at the time of download. This indicates that the downloaded file was not updated to a later version of the data set. Yet, Dataquest gives the information provided by the other columns:

- `Continent_Name`: continent name

- `Two_Letter_Country_Code`: country codes

- `Country_Region`: country names

- `daily_tested`: number of tests conducted on the day; if daily data is unavailable, daily tested is averaged across number of days in between

- `daily_positive`: number of positive cases reported on the day; if daily data is unavailable, daily positive is averaged across number of days in

Raw Data Cleaning

Data Cleaning Q1

The first clean that we do is removing the `Province_State` column. Indeed, it might create some unwanted bias, as it also gives information about the specific province. So, to make the data "nationalized", we avoid looking at precise regions, and consider only those rows with "All States" value:

```
glimpse(covid19_allstates)
```

```
## Rows: 3,781
## Columns: 14
## $ Date                   <chr> "2020-01-20", "2020-01-22", "2020-01-23", "202~
## $ Continent_Name         <chr> "Asia", "North America", "North America", "Asi~
## $ Two_Letter_Country_Code <chr> "KR", "US", "US", "KR", "US", "AU", "GB", "US"~
## $ Country_Region         <chr> "South Korea", "United States", "United States~
## $ Province_State         <chr> "All States", "All States", "All States", "All~
## $ positive               <int> 1, 1, 1, 2, 1, 4, 1, 1, 4, 0, 3, 1, 1, 5, 0, 0~
## $ hospitalized           <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
## $ recovered              <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
## $ death                  <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
## $ total_tested           <int> 4, 1, 1, 27, 1, 0, 31, 1, 0, 3, 51, 52, 1, 0, ~
## $ active                 <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
## $ hospitalizedCurr       <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
## $ daily_tested           <int> 0, 0, 0, 5, 0, 0, 0, 0, 0, 0, 12, 21, 0, 0, 0,~
## $ daily_positive         <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0~
```

As we can see we're left with just 30% of the data we previously had.

The second thing that we need to be aware of is the "unit of measurement" adopted by each column. For example, there are factors with cumulative information, others with daily information. To better address our research questions, we only consider the columns providing daily data, that is `Date`, `Country_Region`, `active`, `hospitalizedCurr`, `daily_tested`, `daily_positive`. We get the following refactored data frame:

```
glimpse(covid19_allstates_daily)
```

```
## Rows: 3,781
## Columns: 6
## $ Date             <chr> "2020-01-20", "2020-01-22", "2020-01-23", "2020-01-24~
## $ Country_Region   <chr> "South Korea", "United States", "United States", "Sou~
## $ active           <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,~
## $ hospitalizedCurr <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,~
## $ daily_tested     <int> 0, 0, 0, 5, 0, 0, 0, 0, 0, 0, 12, 21, 0, 0, 0, 1, 10,~
## $ daily_positive   <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1,~
```

Now data has been cleaned for Q1.

Data Cleaning for Q2

Data Cleaning for Q3

Analysis

Data Analysis for Q1

Recall that with Q1 we are trying to answer the following question: *which countries have had the highest number of positive cases against the number of tests?*

Based on our cleaned data set, we notice that our data is collected on a daily basis. Therefore we can find the ratios on the overall number of positive cases over the total number of tests performed in each country each day. What We want to do is create a data set that groups `Country_Region` and aggregates data over all tests made, positive, active and hospitalized cases in the period of time that the data set covers (i.e. from January 20th to June 1st 2020).

Displaying the countries based on those that have done more tests, we get:

```
## # A tibble: 15 x 5
##    Country_Region    tested positive   active hospitalized
##    <chr>              <int>    <int>    <int>        <int>
##  1 United States   17282363  1877179        0            0
##  2 Russia          10542266   406368  6924890            0
##  3 Italy            4091291   251710  6202214      1699003
##  4 India            3692851    60959        0            0
##  5 Turkey           2031192   163941  2980960            0
##  6 Canada           1654779    90873    56454            0
##  7 United Kingdom   1473672   166909        0            0
##  8 Australia        1252900     7200   134586         6655
##  9 Peru              976790    59497        0            0
## 10 Poland            928256    23987   538203            0
## 11 South Korea       916276    11493   302633            0
## 12 Israel            546626    16647        0        30653
## 13 Germany           518647    29943        0            0
## 14 Belgium           511055    54209   220744            0
## 15 Czech Republic    446758     9321        0            0
```

Now, the result can be easily found by dividing the `positive` cases over the `tested` column and rank the result for each country. We are going to do it using vectors, extracted from the data frame, as an exercise.

We first extract a vector of the 10 countries that have done more tests, over which we are going to concude the rest of the analysis:

```
countries
```

```
## [1] "United States"  "Russia"         "Italy"          "India"
## [5] "Turkey"         "Canada"         "United Kingdom" "Australia"
## [9] "Peru"           "Poland"
```

Note: if we extract the vector on the whole cleaned data set we get that the length of the vector `countries` is 108, which means that we cover more than half of the world with our data. Yet we decide to restrict the study-case on just the top 10 is because we have more data on the tests made, and hence a better chance of having the correct data. As a matter of fact, there might be countries with higher ratios, but more incorrect and non sufficient data, making our data exploration biased (or at least not correct enough).

Then, we extract our two final vectors to run the analysis and assign to each key the name of the country. Thus, we get:

```
tested_cases
```

```
##  United States        Russia         Italy         India        Turkey
##       17282363      10542266       4091291       3692851       2031192
##         Canada United Kingdom     Australia          Peru        Poland
##        1654779       1473672       1252900        976790        928256
```

```
positive_cases
```

```
##  United States        Russia         Italy         India        Turkey
##        1877179        406368        251710         60959        163941
##         Canada United Kingdom     Australia          Peru        Poland
##          90873        166909          7200         59497         23987
```

The conclusion lies in a simple division: `positive_cases / tested_cases`, which yields the following result:

`positive_tested_ratio`

```
##      Australia          India         Poland         Russia         Canada
##    0.005746668    0.016507300    0.025840932    0.038546552    0.054915490
##           Peru          Italy         Turkey  United States United Kingdom
##    0.060910738    0.061523368    0.080711720    0.108618191    0.113260617
```

Data Analysis for Q2

Data Analysis for Q3

Conclusions

This paragraph concisely summarizes the findings of each research question:

- **Q1**: which countries have had the highest number of positive cases against the number of tests?

The following is the final vector we found in the Data Analysis for Q1 section:

```
##      Australia          India         Poland         Russia         Canada
##    0.005746668    0.016507300    0.025840932    0.038546552    0.054915490
##           Peru          Italy         Turkey  United States United Kingdom
##    0.060910738    0.061523368    0.080711720    0.108618191    0.113260617
```

Since there is no precise meaning in the word *highest* that is in question, we decide to choose 4 countries with the highest positive/tests ratio. By doing so, we get:

```
## United Kingdom  United States         Turkey          Italy
##     0.11326062     0.10861819     0.08071172     0.06152337
```

Therefore the top 4 countries were (in order):

1. United Kingdom with a ratio of `0.11326062`

2. United States with a ratio of `0.10861819`

3. Turkey with a ratio of `0.08071172`

4. Italy with a ratio of `0.06152337`

Their corresponding data is displayed in the next matrix:

```
##          Positive/Tested %   Tested Positive
## UK              11.326062  1473672   166909
## US              10.861819 17282363  1877179
## Turkey           8.071172  2031192   163941
## Italy            6.152337  4091291   251710
```

- **Q2**: which countries have made the best effort in terms of the number of COVID-19 tests conducted related to their population?

- **Q3**: which countries were ultimately the most and least affected related to their population?

Improvements

- **Q1**. It was solely based on the countries that had more data on the number of tests made. This is correct but does not give the best result (explain!)

- **Q2**.

- **Q3**.